



E3 — Bioq Throughput Scaling on Aliyun Function Compute


(a) Makespan vs N




(b) Speedup vs N




(c) Throughput vs N




(d) Peak Concurrency vs N



(e) Per-Job Latency Distribution



(f) Cold-Start Overhead (submit → running)



—●— dockq (cold (CPU)) —■— plip (cold (CPU)) —●— reinvent (warm (GPU))
—◆— mmseqs2 (hot (GPU)) —▲— proteinmpnn (hot (GPU)) —▼— rfdiffusion2 (warm (GPU))